

Classic CS

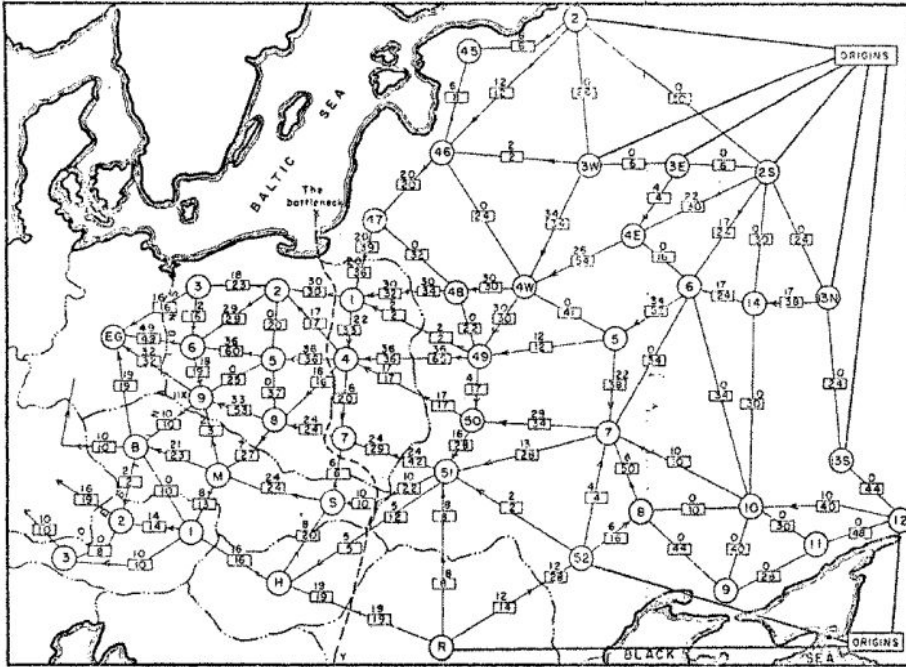
Maximum Flow

Mai Dahshan

updated November 12, 2024



Motivation




Soviet railway network, 1940

Find the **maximum** amount of cargo that can be **transported** from **sources** in the Western Soviet Union to **destinations** in Eastern Europe countries


Flow Network

- A flow network is a directed graph G
- Edges represent pipes that carry flow
- Each edge $\langle u, v \rangle$ has a maximum capacity $c_{\langle u, v \rangle}$
- A source node s in which flow arrives
- A sink node t out which flow leaves



Flow Network


- The flow network problem is defined as follows:
 - Given a directed graph G with non-negative integer weights
 - Each edge stands for the capacity of that edge
 - Two different vertices, s and t , called the source and the sink
 - The source only has out-edges and the sink only has in-edges
 - Find the maximum amount of some commodity that can flow through the network from source to sink



Each edge stands for the capacity of that edge.

Flow Network

- One way to imagine the situation is imagining each edge as a pipe that allows a certain flow of a liquid
 - The source is where the liquid is pouring from, and the sink is where it ends up.
 - Each edge weight specifies the maximal amount of liquid that can flow through that pipe per second.
 - Given that information, what is the most liquid that can flow from source to sink in the steady state?



Flow Network Examples



- **Transportation:** Modeling traffic on a network of roads, or the routing of packages by a company
- **Communication:** Routing packets in a communication network
Air travel: Sequencing the legs of a flight
- **Railway systems:** Transporting goods across a railway system
Vehicle routing: Finding the best routes for delivery trucks to minimize costs and time

Flow Graph

- **Flow graph** is a directed graph with distinguished vertices s (source) and t (sink)
- Capacities on the edges, $c(e) \geq 0$
- Assign flows $f(e)$ to the edges such that:
 - **Capacity Rule:** $0 \leq f(e) \leq c(e)$
 - **Conservation Rule:** Flow is conserved at vertices other than s and t
 - Flow conservation: flow going into a vertex equals the flow going out
 - The flow leaving the source is as large as possible


Flow Graph

- A **flow** in a flow network is function f that assigns each edge e a non-negative integer value, namely the flow.
- The function has to fulfill the following two conditions:
 - $0 \leq f(e) \leq c(e)$
 - The **sum of the incoming flow** of a vertex u has to be **equal to the sum of the outgoing flow** of except in the source and sink vertices



Flow Graph



The ***flow of the network*** is defined as the flow from the source, or into the sink



The network flow is 23



Greedy Approach

- In a **greedy algorithm**, at each step, you would choose the edge that appears to maximize flow from the source to the sink, based on a local criterion (e.g., choosing the edge with the maximum available capacity)





Greedy Approach

- Start with $f(e) = 0$ for all edge $e \in E$
- Find an s-t path P where each edge maximize $f(e)$ and $f(e) < c(e)$
- Augment flow along path P
- Repeat until you get **stuck**




Greedy Approach




locally optimality != global optimality


Greedy Approach



Greedy Approach




Step 1:



Greedy Approach


Step 2:



Flow = 10 + 1 = 11

Greedy = 11

Max Flow:



Optimal = 12

Paths {s,a,c,t}, {s,a,d,t}, {s,b,c,t}, and {s,b,d,t}

Ford-Fulkerson Algorithm



- The **Ford-Fulkerson algorithm** is a classic method for finding the **maximum flow** in a flow network. It uses the concept of **augmenting paths** and operates on the **residual graph** to iteratively improve the flow until no more augmenting paths exist

Residual graph

- Residual graph shows the remaining capacity
- It tracks the flow capacity and allows for flow adjustments during the process of finding maximum flow.
- Given a flow f in graph G , the residual graph G_f models additional flow that is possible
 - **Forward edge** for each edge in G with weights set to remaining capacity $c(e)-f(e)$
 - models **additional flow** that can be sent along edge
 - **Backward edge** by flipping each edge e in G with weight set to the flow $f(e)$
models amount of flow that can be **removed** from the edge


Residual graph

- **Forward edge capacity** corresponds to the amount of flow that can still be pushed through that edge ($c(e)-f(e)$)
- A **forward edge exists** if there is any **remaining capacity** for flow in the original edge.
- If the **flow on a particular edge reaches its capacity**, then the forward edge in the residual graph will have **zero residual capacity**.
- If **no flow** has been sent along the edge, the **forward edge will have the same capacity as the original edge** in the flow network




Residual graph

- **Backward Edge** represents the possibility of "undoing" or "reversing" the flow along a previously used edge
- It exists when there is flow already passing from u to v , and the flow can be "pushed back" from v to u . This backward edge allows us to adjust the flow if necessary
- The **capacity** of the **backward edge in the residual graph** is the **amount of flow that has already been sent along the forward edge in the original graph**
- The backward edge allows the algorithm to adjust the flow when a previously used path is no longer optimal




Flow graph




Residual graph

Ford-Fulkerson Algorithm




Flow graph G




Residual Graph G_f

Ford-Fulkerson Algorithm



Flow graph G




Residual Graph G_f

Ford-Fulkerson Algorithm


- In a residual graph, an **augmented path** is a path from the source node s to the sink node t where each edge has **positive residual capacity** (either forward or backward). The path can mix both types of edges, as follows:

- **Forward edges** (with positive residual capacity): These edges allow for additional flow to be pushed from the source to the sink
- **Backward edges** (with positive residual capacity): These edges allow the possibility of undoing or adjusting flow by pushing flow in the reverse direction




Ford-Fulkerson Algorithm

- Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight
- Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$
- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)



Flow graph G




Residual Graph G_f

Ford-Fulkerson Algorithm


Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight

Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$

- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)
- Remove $w(e)$ flow along all **backward** edges (these contain at least $w(e)$ units of flow)



Flow graph G




Residual Graph G_f

Ford-Fulkerson Algorithm


Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight
Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$

- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)
- Remove $w(e)$ flow along all **backward** edges (these contain at least $w(e)$ units of flow)

Observe: Flow has increased by $w(e)$



Flow graph G



Residual Graph G_f

Ford-Fulkerson Algorithm

Define an augmenting path to be an $s \rightarrow t$ path in the residual graph G_f (using edges of non-zero weight)


Ford-Fulkerson max-flow algorithm:

- Initialize $f(e) = 0$ for all $e \in E$
- Construct the residual network G_f
- While there is an augmenting path p in G_f :
 - Let $c = \min_{e \in p} c_f(e)$ ($c_f(e)$ is the weight of edge e in the residual network G_f)
 - Add c units of flow to G based on the augmenting path p
 - Update the residual network G_f for the updated flow


Ford-Fulkerson approach:
take any augmenting path

Ford-Fulkerson Algorithm- Example 1

- Initially set the flow along every edge to 0
- Construct a residual graph for this network. It should look the same as the input flow network




Flow graph



Residual graph


Ford-Fulkerson Algorithm- Example 1

- Use a pathfinding algorithm like (DFS) or (BFS) to find a path P from s to t that has available capacity **in the residual graph**



Ford-Fulkerson Algorithm- Example 1


- Let $cap(P)$ indicate the maximum amount of stuff that can flow along this path
 - To find the capacity of this path, we need to look at all edges e on the path and subtract their current flow, from their capacity. We'll set $cap(P)$ to be equal to the smallest value since this will **bottleneck the path**



3 is the bottleneck for this the path



Ford-Fulkerson Algorithm- Example 1

- We then **augment the flow** across the forward edges in the path P by adding $cap(P)$ value. For flow across the back edges in the residual graph, we subtract our $cap(P)$ value
- Update the residual graph with these flow adjustments





Ford-Fulkerson Algorithm- Example 1

- Search through the updated residual graph for a new s - t path. There are no forward edges available anymore, but we can use a back edge to augment the current flow




Ford-Fulkerson Algorithm- Example 1


- There are now no edges with available capacity that we can use to create a path from s to t . This means our run of the Ford-Fulkerson algorithm is complete and our max flow leading into t is 5



Ford-Fulkerson Algorithm- Example 2





Initially: $f(e) = 0$ for all $e \in E$





Residual graph G_f



Ford-Fulkerson Algorithm- Example 2





Ford-Fulkerson Algorithm- Example 2





Ford-Fulkerson Algorithm- Example 2





Ford-Fulkerson Algorithm- Example 2





Ford-Fulkerson Algorithm- Example 2



Ford-Fulkerson Algorithm- Example 2

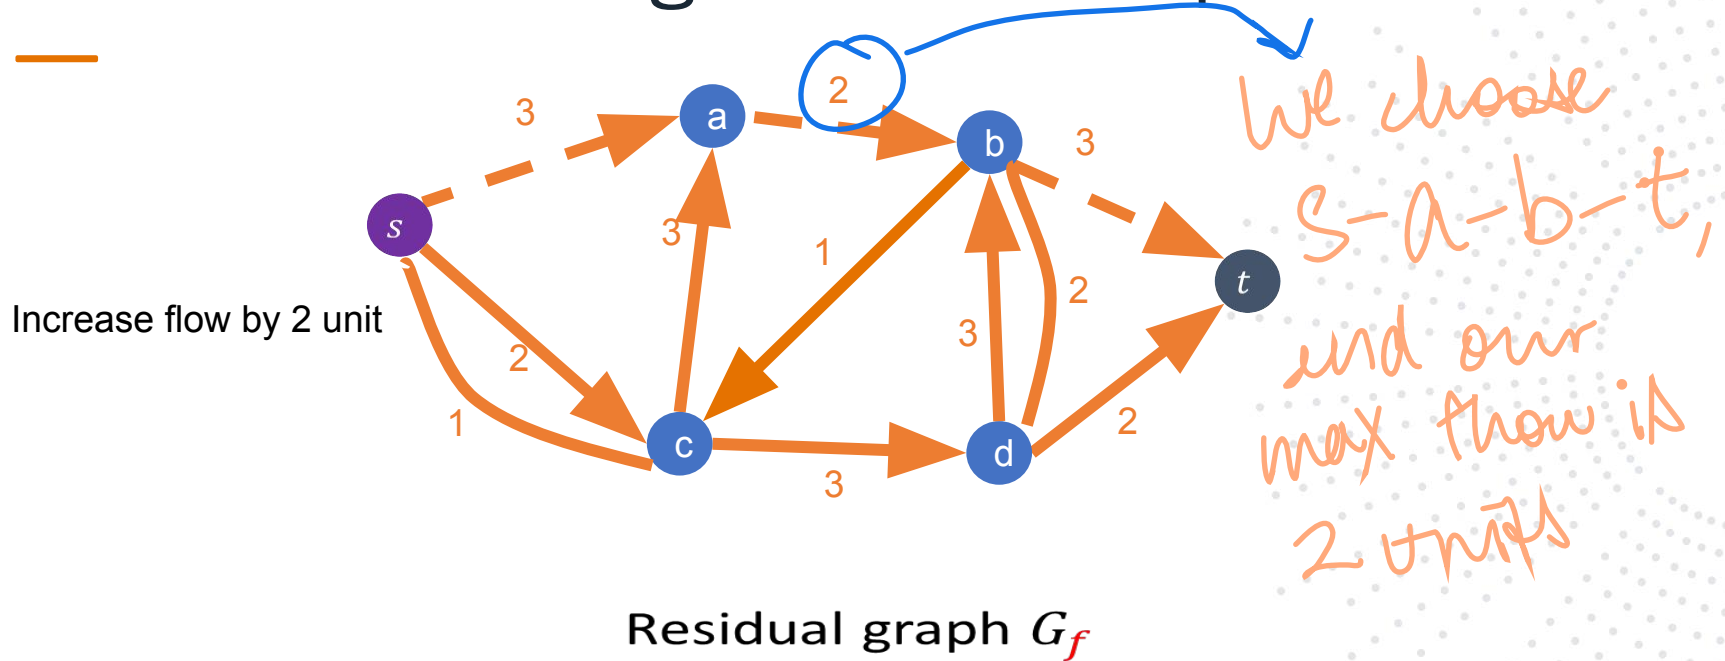


Ford-Fulkerson Algorithm- Example 3 Sol.1

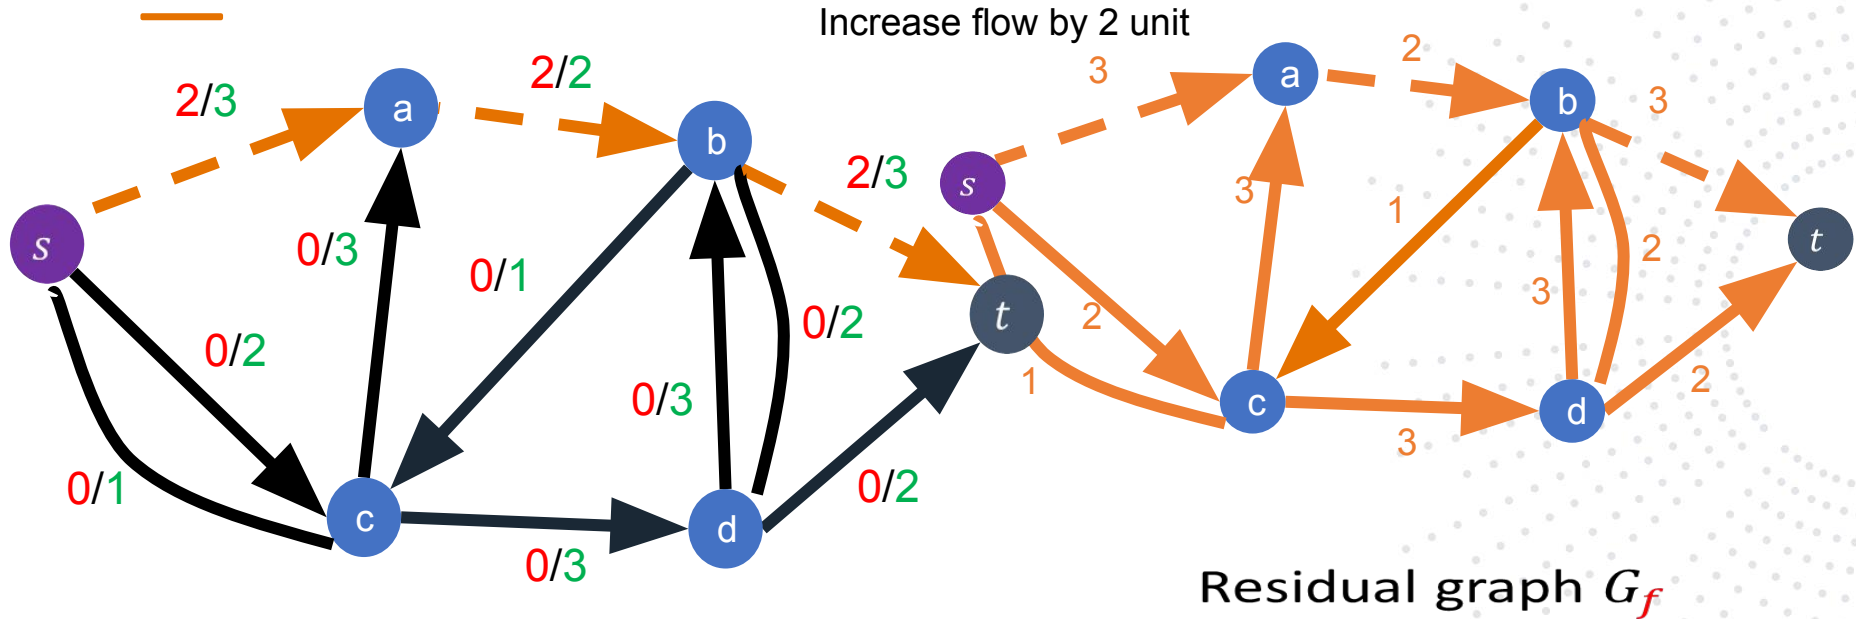


Initially: $f(e) = 0$ for all $e \in E$

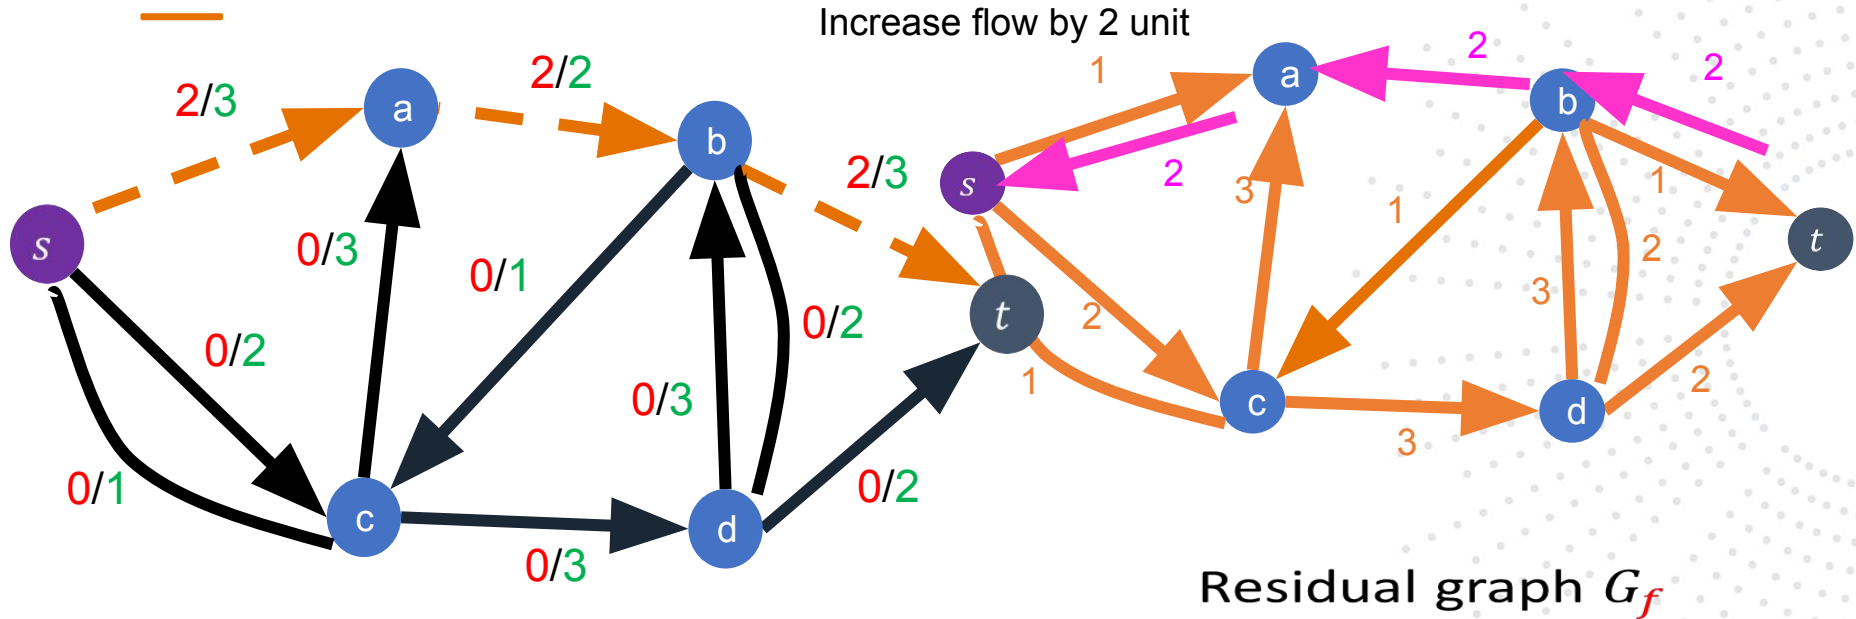
Ford-Fulkerson Algorithm- Example 3 Sol.1



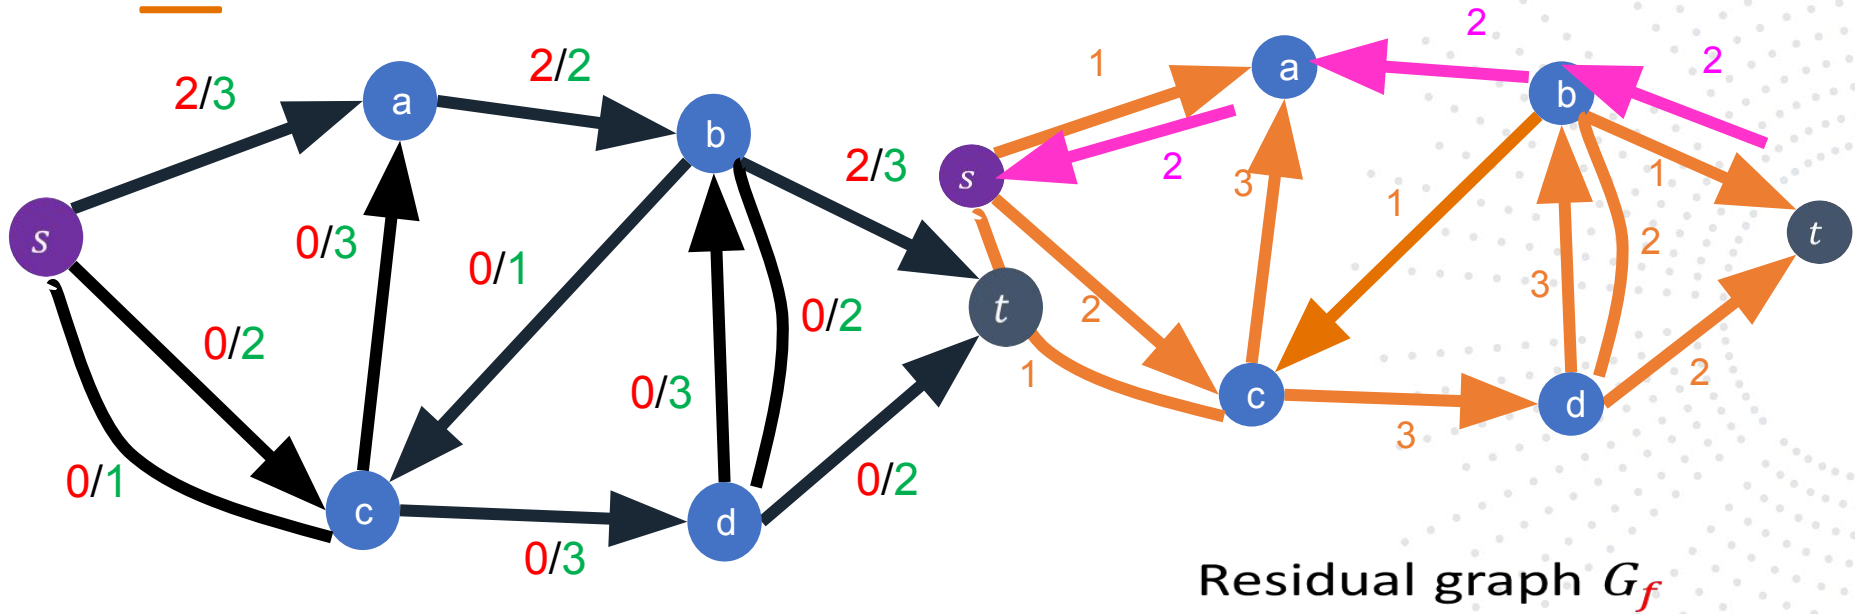
Ford-Fulkerson Algorithm- Example 3 Sol.1



Ford-Fulkerson Algorithm- Example 3 Sol.1




Ford-Fulkerson Algorithm- Example 3 Sol.1




Ford-Fulkerson Algorithm- Example 3 Sol.1

Increase flow by 2 unit

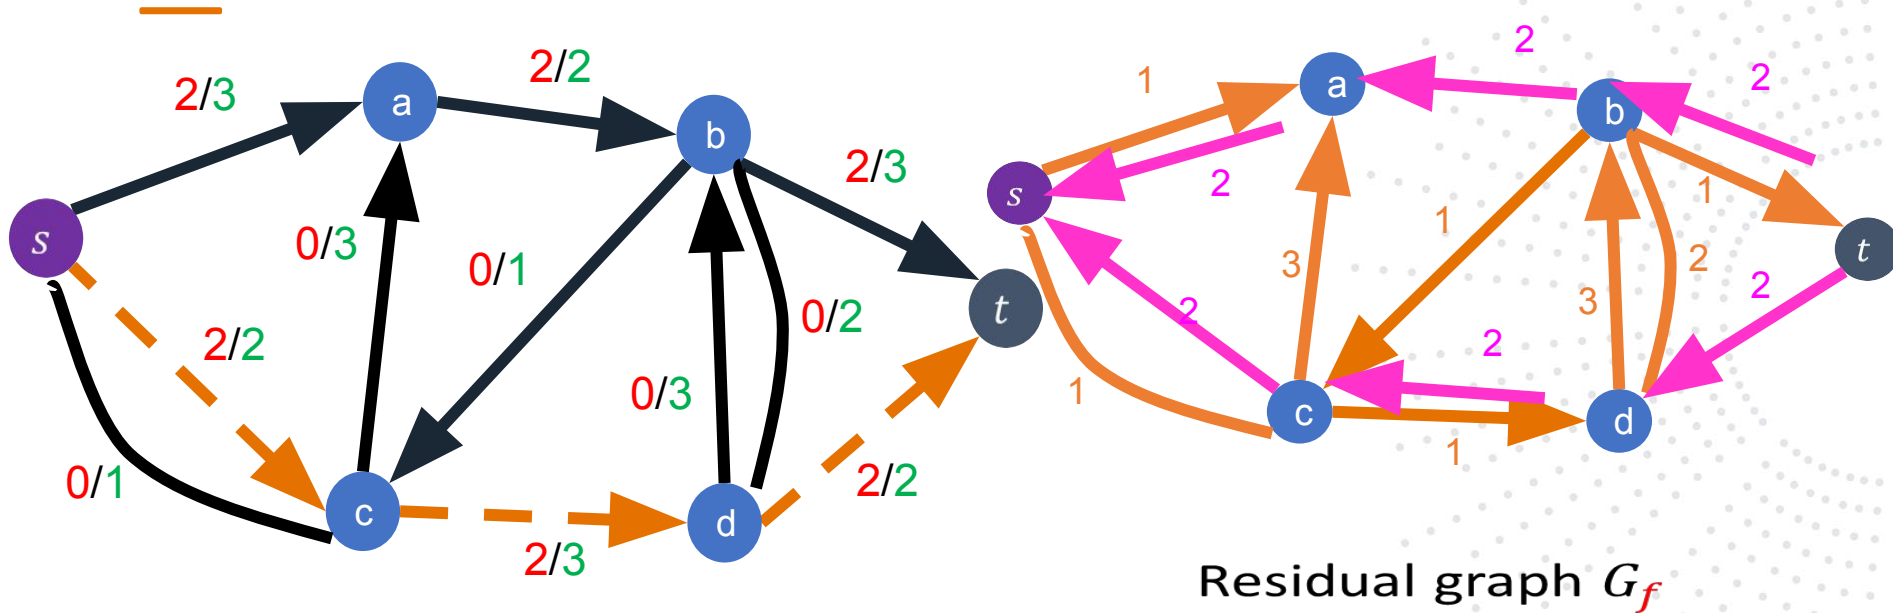


Residual graph G_f



Ford-Fulkerson Algorithm- Example 3 Sol.1



Ford-Fulkerson Algorithm- Example 3 Sol.1



Ford-Fulkerson Algorithm- Example 3 Sol.1





No more augmented paths

Residual graph G_f


Maximum flow: 4

Ford-Fulkerson Algorithm- Example 3 Sol.2




Initially: $f(e) = 0$ for all $e \in E$

Ford-Fulkerson Algorithm- Example 3 Sol.2

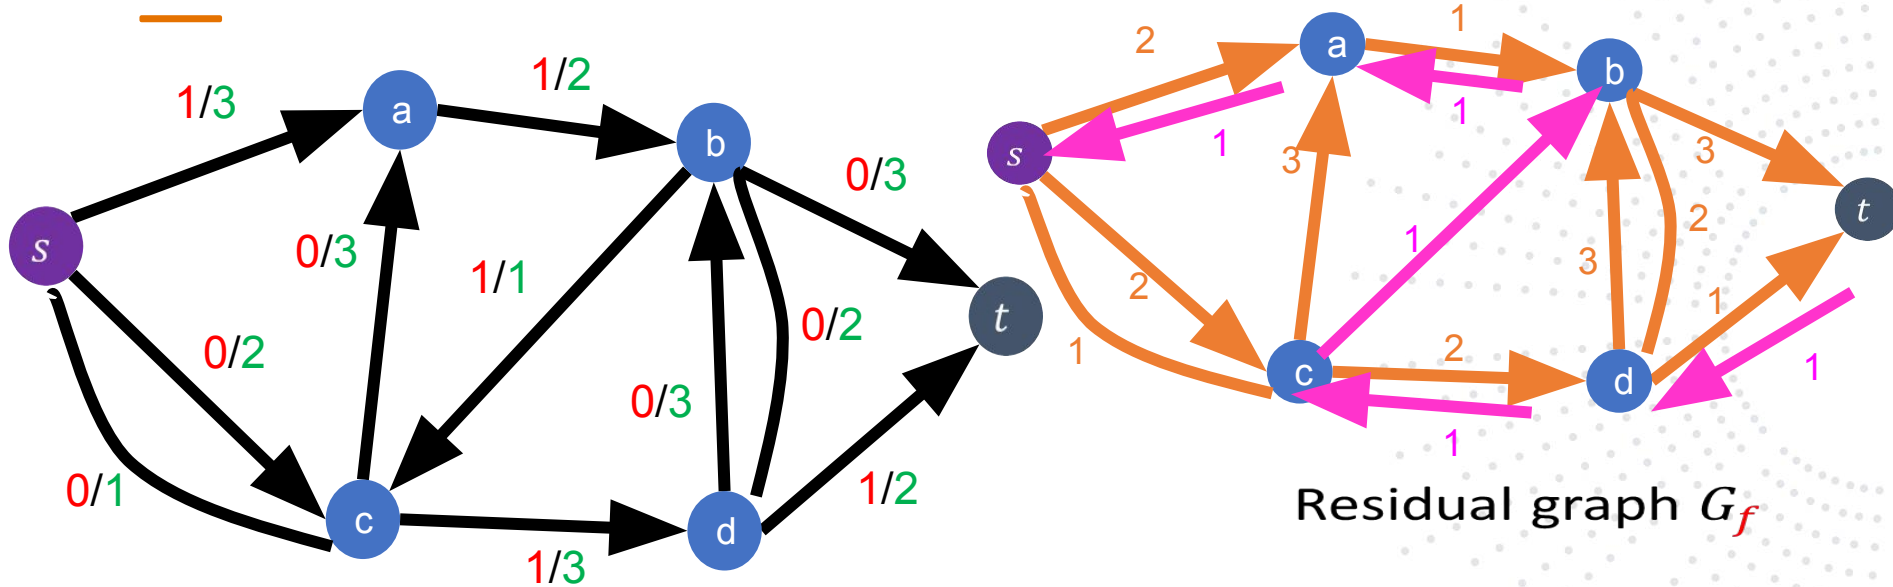


Residual graph G_f


Ford-Fulkerson Algorithm- Example 3 Sol.2




Ford-Fulkerson Algorithm- Example 3 Sol.2





Ford-Fulkerson Algorithm- Example 3 Sol.2



Ford-Fulkerson Algorithm- Example 3 Sol.2




Ford-Fulkerson Algorithm- Example 3 Sol.2




Residual graph G_f


Ford-Fulkerson Algorithm- Example 3 Sol.2



Ford-Fulkerson Algorithm- Example 3 Sol.2




Ford-Fulkerson Algorithm- Example 3 Sol.2




Ford-Fulkerson Algorithm- Example 3 Sol.2

Increase flow by 1 unit




Residual graph G_f

Ford-Fulkerson Algorithm- Example 3 Sol.2



Residual graph G_f

Ford-Fulkerson Algorithm- Example 3 Sol.2




Maximum flow: 4


Worst-Case Ford-Fulkerson Algorithm

- The **worst-case time complexity** of the **Ford-Fulkerson algorithm** can arise in specific scenarios where the algorithm has to process a large number of augmenting paths to reach the maximum flow
 - Path augmentation is slow
 - Edge capacities are small

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm



Flow graph

Increase flow by 1 unit



Residual graph

Ford-Fulkerson algorithm

- **Worst Case (Exponential Time Complexity):**
 - $O(E \cdot \text{Max Flow})$, which can be exponential in some cases
- This happens when the algorithm uses **Depth-First Search (DFS)** to find augmenting paths, and capacities lead to tiny flow increments.
- For example, if each augmenting path adds only 1 unit of flow and the maximum flow is F , there can be F iterations. If E edges are checked in each iteration, the time complexity becomes $O(E \cdot F)$

Edmonds-Karp Algorithm

- The **Edmonds-Karp algorithm** is a specific implementation of the **Ford-Fulkerson method** for solving the **maximum flow problem** in a flow network. It improves the Ford-Fulkerson method by using **Breadth-First Search (BFS)** to find the shortest augmenting path in terms of the number of edges

Edmonds-Karp Algorithm

Ford-Fulkerson max-flow algorithm:

- Initialize $f(e) = 0$ for all $e \in E$
- Construct the residual network G_f
- While there is an augmenting path in G_f , let p be the path with fewest hops:
 - Let $c = \min_{e \in E} c_f(e)$ ($c_f(e)$ is the weight of edge e in the residual network G_f)

How to find this?
Use breadth-first search (BFS)!

Edmonds-Karp = Ford-Fulkerson using
BFS to find augmenting path


Edmonds-Karp Algorithm

- **Edmonds-Karp algorithm** has a time complexity of $O(V \cdot E^2)$
 - BFS takes $O(V+E)$ time per iteration
 - Each augmenting path can increase the flow by a finite amount.
 - The maximum number of BFS iterations is $O(V \cdot E)$, because each edge is involved at most V times.
- Since each BFS takes $O(V+E)$, and BFS is called $O(V \cdot E)$ times, the total time complexity is:


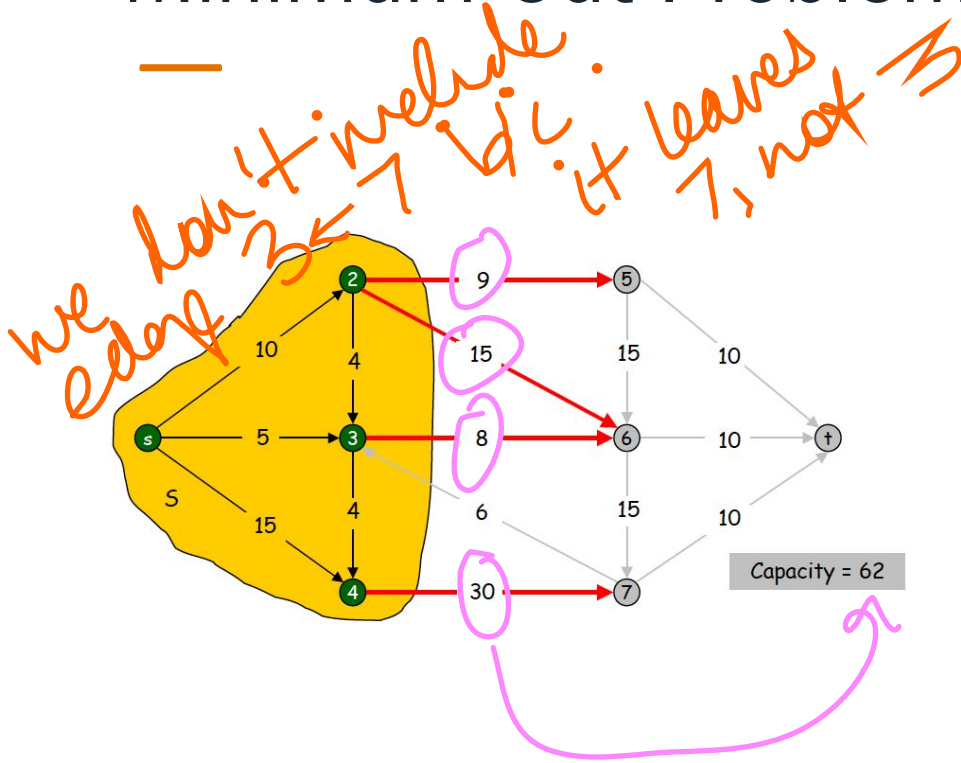
$$O((V+E) \cdot V \cdot E) = O(V \cdot E^2)$$

Minimum Cut Problem

- Given a graph $G=(V,E)$ with vertices V and edges E , and two designated vertices s (source) and t (sink)
- The **min-cut** is the smallest set of edges whose removal disconnects the graph or separates s from t .
- A cut in a graph is a partition of the vertices V into two disjoint subsets S and T
 - s is in S and t is in T
 - $S \cup T = V$ and $S \cap T = \emptyset$
- The **capacity of a cut** is the sum of the weights crossing from S to T
 - $capacity(S, T) = \text{sum of weights of edges leaving } S$




Minimum Cut Problem



Minimum Cut Problem

- An edge crosses the cut if it goes from s's side to t's side




This cut has capacity =>
 $4 + 2 + 10 = 16$

Flows and Cuts


- Let f be a flow, and let (S, T) be any s - t cut. Then, the net flow sent across the cut is equal to the amount reaching t

net flow is different from the cut

- Let f be a flow, and let (S, T) be any s - t cut. Then the value of the flow is at most the capacity of the cut




Cut capacity = 30 Flow value \leq 30




Max Flow and Min Cut

- Let f be a flow, and let (S, T) be an s - t cut whose capacity equals the value of f . Then f is a max flow and (S, T) is a min cut




Max-Flow Min-Cut Theorem

Max-flow min-cut theorem. (Ford-Fulkerson, 1956): In any network, the value of max flow equals capacity of min cut




Max-Flow Min-Cut Theorem

Let f be a flow in a graph G




Augmenting path theorem: A flow f is a max flow if and only if there are no augmenting paths

Max-Flow Min-Cut Theorem



Max flow: 4
Min cut: 4




When there are no more augmenting paths in the graph, there is a **cut** whose cost/capacity matches the **flow**

Why Study Flow Networks?

- Unlike divide-and-conquer, greedy, or DP, flow network doesn't seem like an algorithmic framework
 - It seems more like a single problem
- How this single problem be useful?


Bipartite Matching Problem

- A Bipartite Graph $G = (V, E)$ is a graph in which the vertex set V can be divided into two disjoint subsets X and Y such that every edge $e \in E$ has one endpoint in X and the other end point in Y
- A matching M is a subset of edges such that each node in V appears in at most one edge in M




Bipartite Matching Problem

- Given a set of people L and set of jobs R
- Each person can do only some of the jobs
- How to model this relation as bipartite graph?



Bipartite Matching Problem

- A matching gives an assignment of people to tasks
- We want to get as many tasks done as possible
- So, want a maximum matching: one that contains as many edges as possible

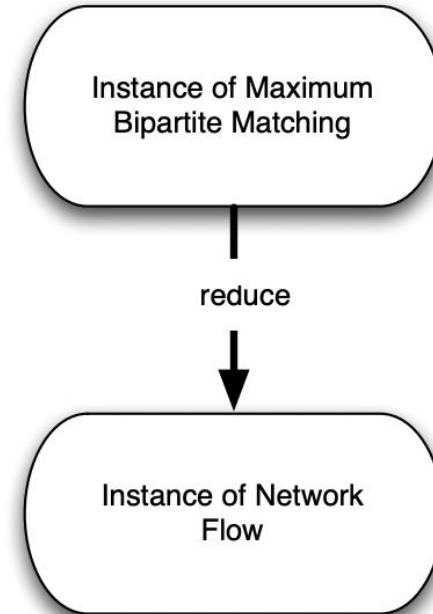


Bipartite Matching Problem

- Our goal is to get the **maximum matching**
- **Maximum matching** is a matching that contains the **maximum number of edges** possible. This is the most "complete" matching in terms of the number of edges that can be added without violating the matching condition


Reduction

- Given an instance of bipartite matching
- Create an instance of network flow
- The solution to the network flow problem can easily be used to find the solution to the bipartite matching



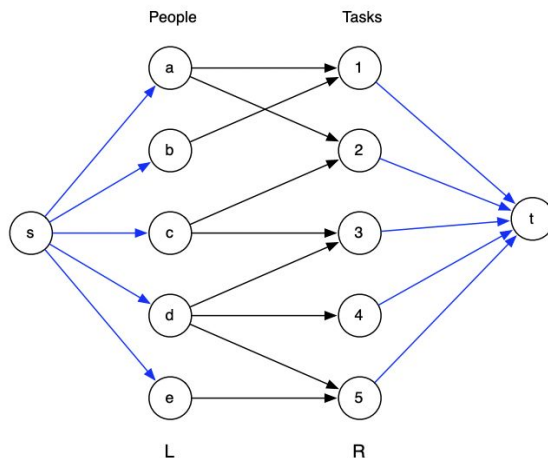
Using Flow Network to Solve Bipartite Matching

1. Given bipartite graph $G = (A \cup B, E)$, direct the edges from A to B



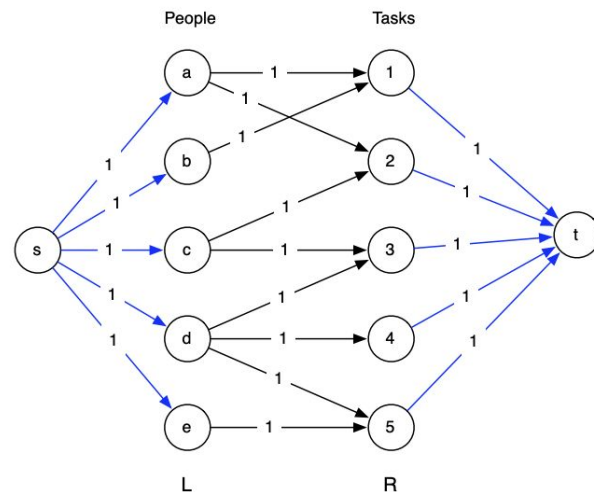
Using Flow Network to Solve Bipartite Matching

1. Given bipartite graph $G = (A \cup B, E)$, direct the edges from A to B
2. Add new vertices s and t
3. Add an edge from s to every vertex in A
4. Add an edge from every vertex in B to t



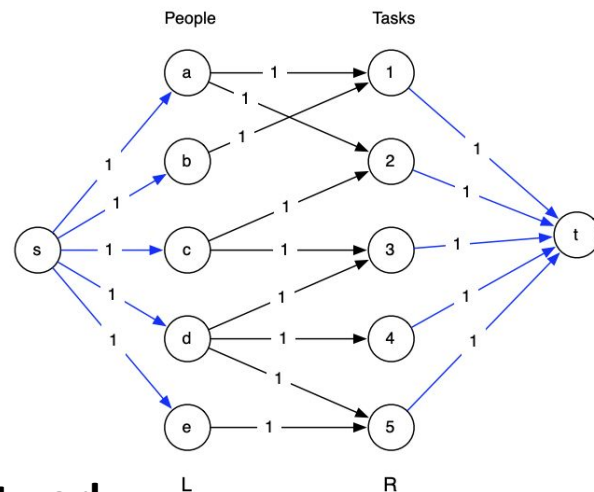
Using Flow Network to Solve Bipartite Matching

1. Given bipartite graph $G = (A \cup B, E)$, direct the edges from A to B
2. Add new vertices s and t
3. Add an edge from s to every vertex in A
4. Add an edge from every vertex in B to t
5. Make all the capacities 1
6. Solve maximum network flow problem on this new graph G



Using Flow Network to Solve Bipartite Matching


1. Given bipartite graph $G = (A \cup B, E)$, direct the edges from A to B
2. Add new vertices s and t
3. Add an edge from s to every vertex in A
4. Add an edge from every vertex in B to t
5. Make all the capacities 1
6. Solve maximum network flow problem on this new graph G



The edges used in the maximum network flow will correspond to the largest possible matching!

Using Flow Network to Solve Bipartite Matching

- We can choose at most one edge leaving any node in A
- We can choose at most one edge entering any node in B



If we chose more than 1, we will **violate the conservation rule** in flow network