

Classic CS

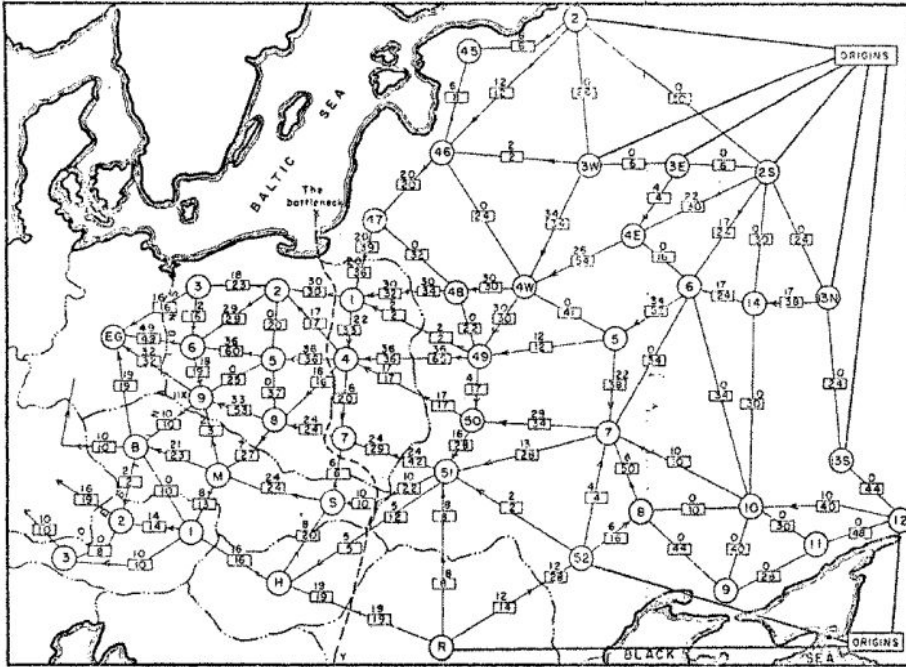
Maximum Flow

Mai Dahshan

November 11, 2024



Motivation




Soviet railway network, 1940

Find the **maximum** amount of cargo that can be **transported** from **sources** in the Western Soviet Union to **destinations** in Eastern Europe countries

Flow Network


- A flow network is a directed graph G
- Edges represent pipes that carry flow
- Each edge $\langle u, v \rangle$ has a maximum capacity $c_{\langle u, v \rangle}$
- A source node s in which flow arrives
- A sink node t out which flow leaves



Flow Network

- The flow network problem is defined as follows:
 - Given a directed graph G with non-negative integer weights
 - Each edge stands for the capacity of that edge
 - Two different vertices, s and t , called the source and the sink
 - The source only has out-edges and the sink only has in-edges
 - Find the maximum amount of some commodity that can flow through the network from source to sink


from $s-a-b-t$,
the max you can
transfer is the min
edge



Each edge stands for the capacity of that edge.

Flow Network

- One way to imagine the situation is imagining each edge as a pipe that allows a certain flow of a liquid
 - The source is where the liquid is pouring from, and the sink is where it ends up.
 - Each edge weight specifies the maximal amount of liquid that can flow through that pipe per second.
 - Given that information, what is the most liquid that can flow from source to sink in the steady state?



Flow Network Examples



- **Transportation:** Modeling traffic on a network of roads, or the routing of packages by a company
- **Communication:** Routing packets in a communication network
Air travel: Sequencing the legs of a flight
- **Railway systems:** Transporting goods across a railway system
Vehicle routing: Finding the best routes for delivery trucks to minimize costs and time

Flow Graph

- **Flow graph** is a directed graph with distinguished vertices s (source) and t (sink)
- Capacities on the edges, $c(e) \geq 0$
- Assign flows $f(e)$ to the edges such that:
 - **Capacity Rule:** $0 \leq f(e) \leq c(e)$
 - **Conservation Rule:** Flow is conserved at vertices other than s and t
 - Flow conservation: flow going into a vertex equals the flow going out
 - The flow leaving the source is as large as possible

Flow Graph

The **flow of the network** is defined as the flow from the source, or into the sink.



$$4 + 7 + 12 = 23$$

The network flow is 23

max Flow from s-a-b-t is 12

min Flow from s-c-d-t is 4

Whatever comes in or out of the node needs to be the same

Flow Graph





Of all valid flows through the graph, find the one that maximizes:

$$|f| = \text{outflow}(s) - \text{inflow}(s)$$



Greedy Approach

- Start with $f(e) = 0$ for all edge $e \in E$
- Find an s-t path P where each edge has $\max f(e)$, where $f(e) < c(e)$
- Augment flow along path P
- Repeat until you get stuck





Greedy Approach

- Start with $f(e) = 0$ for all edge $e \in E$
- Find an s-t path P where each edge has $\max f(e)$, where $f(e) < c(e)$
- Augment flow along path P
- Repeat until you get **stuck**




Greedy Approach



locally optimality != global optimality

Greedy looks for local optimality, not global optimality

Greedy Approach




Using greedy
path 1 of 2 $S-a-c-t \{10$
path 2 of 2 $S-b-d-t \{1$
Sink


note: the paths
have to be unique.

If I wouldn't choose path
 $S-b-c-t$ b/c $c-t$
was already taken in $S-a-c-t$

Greedy Approach




Step 1:




Greedy Approach

Step 2:



Flow = $10 + 1 = 11$

Greedy = 11



Ford-Fulkerson Algorithm


- The **Ford-Fulkerson algorithm** is a classic method for finding the **maximum flow** in a flow network. It uses the concept of **augmenting paths** and operates on the **residual graph** to iteratively improve the flow until no more augmenting paths exist

Ford-Fulkerson Algorithm


Residual Graphs have
forward & backward arrows

- Residual graph shows the remaining capacity
- Given a flow f in graph G , the residual graph G_f models additional flow that is possible
 - **Forward edge** for each edge in G with weights set to remaining capacity $c(e)-f(e)$ forward: how many unused arrows are there?
 - models **additional flow** that can be sent along edge
 - **Backward edge** by flipping each edge e in G with weight set to the flow $f(e)$ models amount of flow that can be **removed** from the edge
backward: how many points can I remove from the graph

Ford-Fulkerson Algorithm




Flow graph G




Residual Graph G_f

Forward: $c(e) - f(e)$

Ford-Fulkerson Algorithm



Flow graph G




Residual Graph G_f

The residual graph tries to help you determine what is the max capacity in your graph


Ford-Fulkerson Algorithm

- Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight
- Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$
- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)



Flow graph G

augmented paths in residual graph (dashed arrows)




Residual Graph G_f

Ford-Fulkerson Algorithm


Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight

Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$

- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)
- Remove $w(e)$ flow along all **backward** edges (these contain at least $w(e)$ units of flow)



Flow graph G




Residual Graph G_f

Ford-Fulkerson Algorithm


Consider a path from $s \rightarrow t$ in G_f using only edges with positive (non-zero) weight
Consider the minimum-weight edge e along the path: we can increase the flow by $w(e)$

- Send $w(e)$ flow along all **forward** edges (these have at least $w(e)$ capacity)
- Remove $w(e)$ flow along all **backward** edges (these contain at least $w(e)$ units of flow)

Observe: Flow has increased by $w(e)$



Flow graph G



Residual Graph G_f

Ford-Fulkerson Algorithm

Define an augmenting path to be an $s \rightarrow t$ path in the residual graph G_f (using edges of non-zero weight)


Ford-Fulkerson max-flow algorithm:

- Initialize $f(e) = 0$ for all $e \in E$
- Construct the residual network G_f
- While there is an augmenting path p in G_f :
 - Let $c = \min_{e \in p} c_f(e)$ ($c_f(e)$ is the weight of edge e in the residual network G_f)
 - Add c units of flow to G based on the augmenting path p
 - Update the residual network G_f for the updated flow


Ford-Fulkerson approach:
take any augmenting path

Ford-Fulkerson Algorithm- Example

- Initially set the flow along every edge to 0
- Construct a residual graph for this network. It should look the same as the input flow network



Flow graph




Residual graph

only forward edges

min flow
is 3


Ford-Fulkerson Algorithm- Example

- Use a pathfinding algorithm like (DFS) or (BFS) to find a path P from s to t that has available capacity **in the residual graph**



Ford-Fulkerson Algorithm- Example


- Let $cap(P)$ indicate the maximum amount of stuff that can flow along this path
 - To find the capacity of this path, we need to look at all edges e on the path and subtract their current flow, from their capacity. We'll set $cap(P)$ to be equal to the smallest value since this will **bottleneck the path**



3 is the bottleneck for this the path


Ford-Fulkerson Algorithm- Example

- We then **augment the flow** across the forward edges in the path P by adding $cap(P)$ value. For flow across the back edges in the residual graph, we subtract our $cap(P)$ value
- Update the residual graph with these flow adjustments




Ford-Fulkerson Algorithm- Example

- Search through the updated residual graph for a new s-t path. There are no forward edges available anymore, but we can use a back edge to augment the current flow

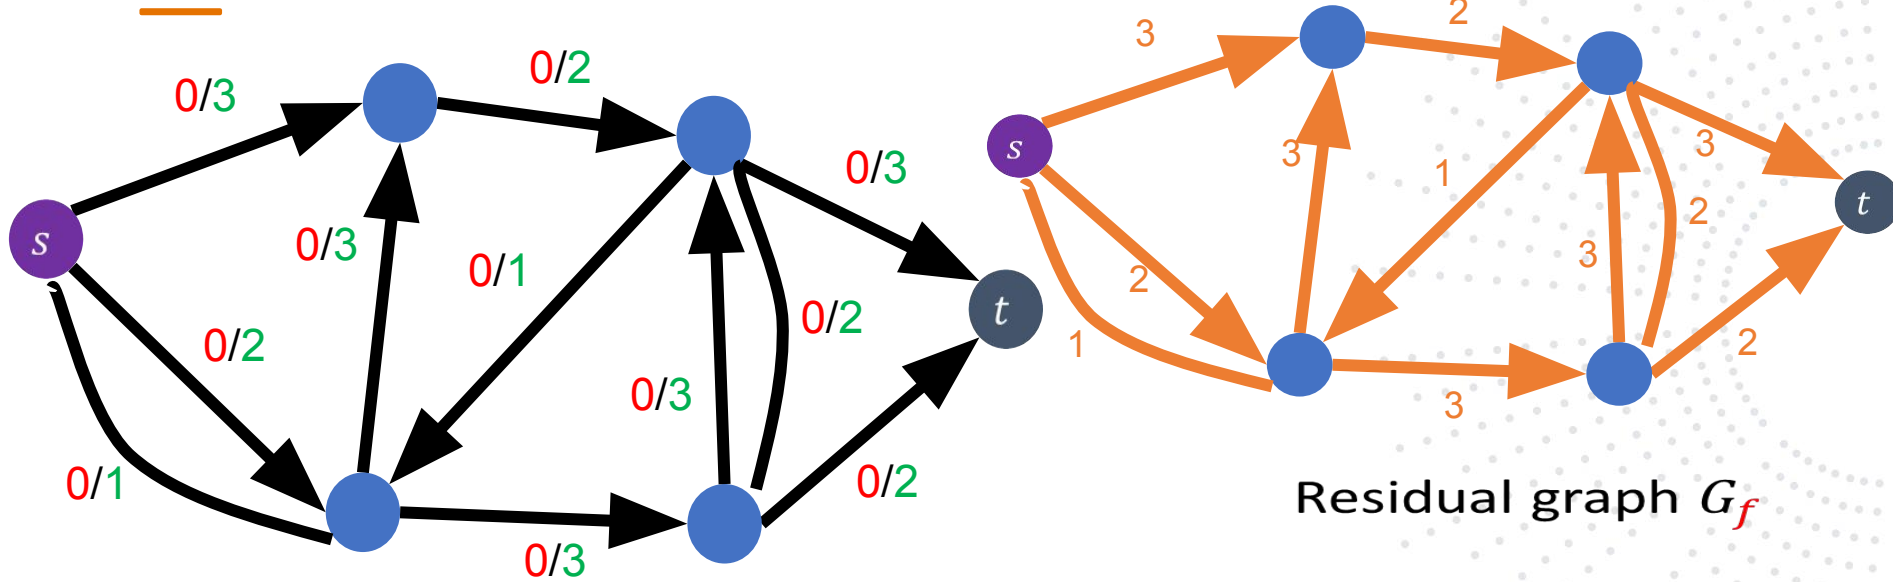


Ford-Fulkerson Algorithm- Example

- There are now no edges with available capacity that we can use to create a path from s to t . This means our run of the Ford-Fulkerson algorithm is complete and our max flow leading into t is 5




Ford-Fulkerson Algorithm- Example




Initially: $f(e) = 0$ for all $e \in E$

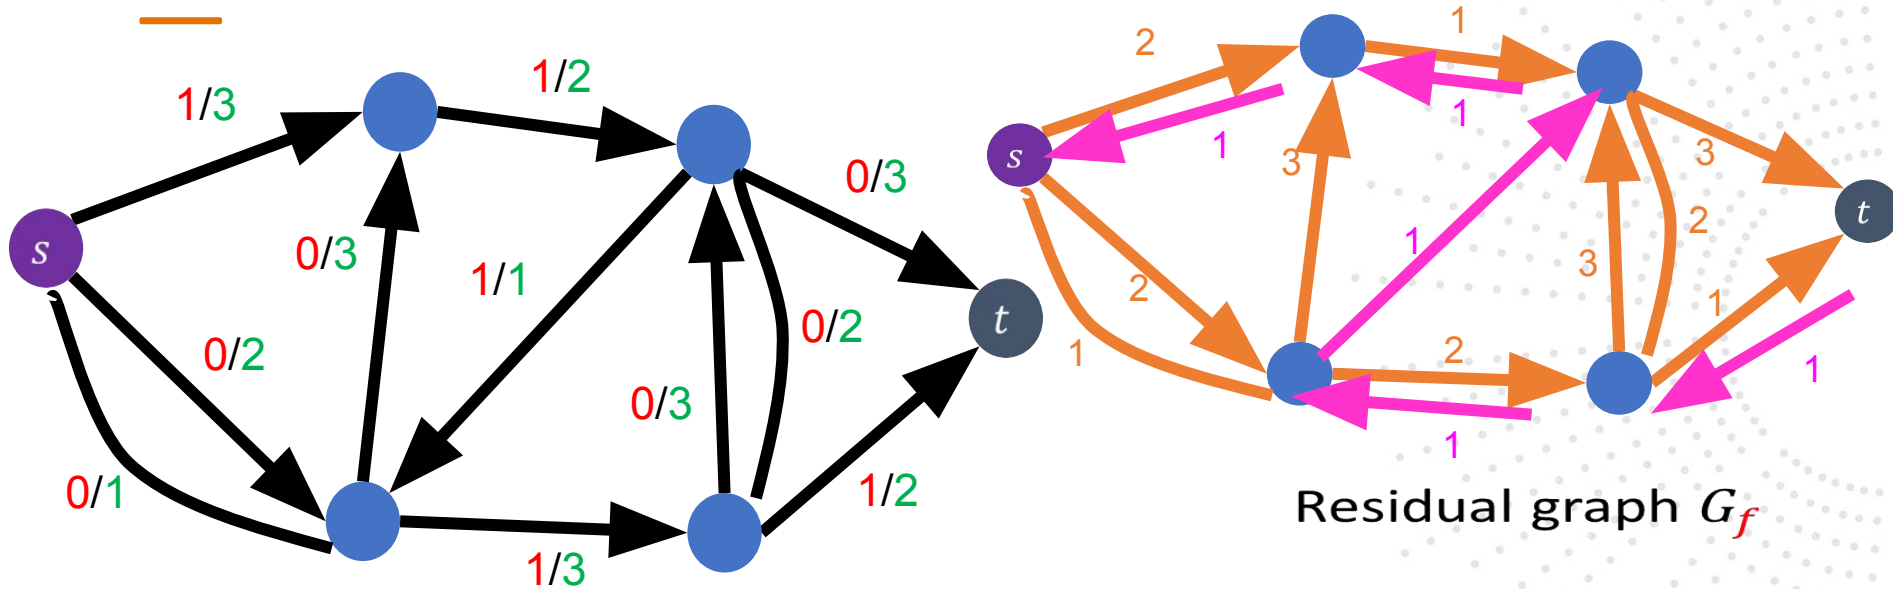
Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm- Example




Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm- Example


—

Increase flow by 1 unit





Residual graph G_f

Ford-Fulkerson Algorithm- Example




Ford-Fulkerson Algorithm- Example




Ford-Fulkerson Algorithm- Example

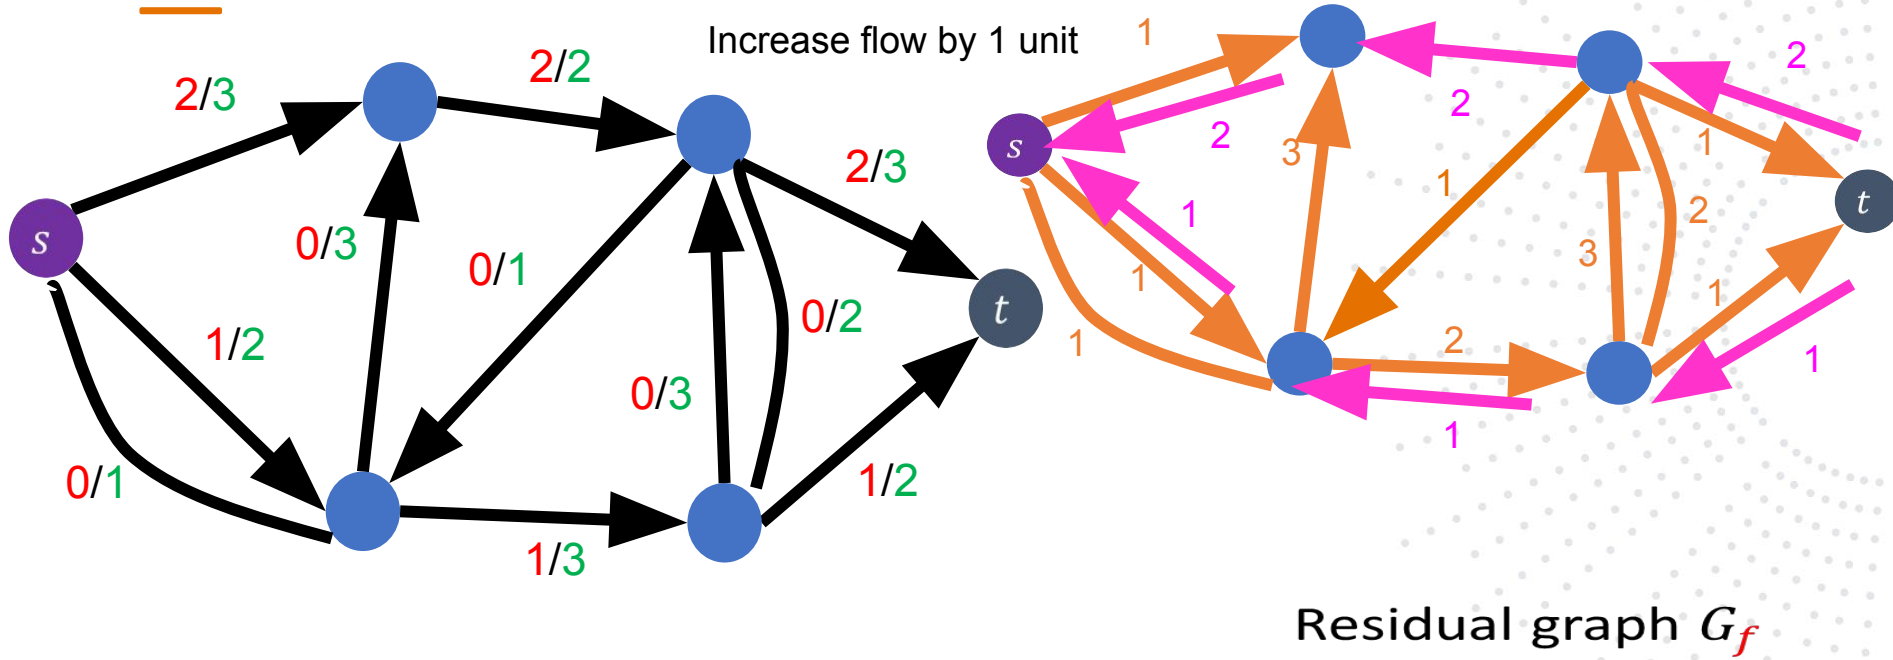
—




Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm- Example

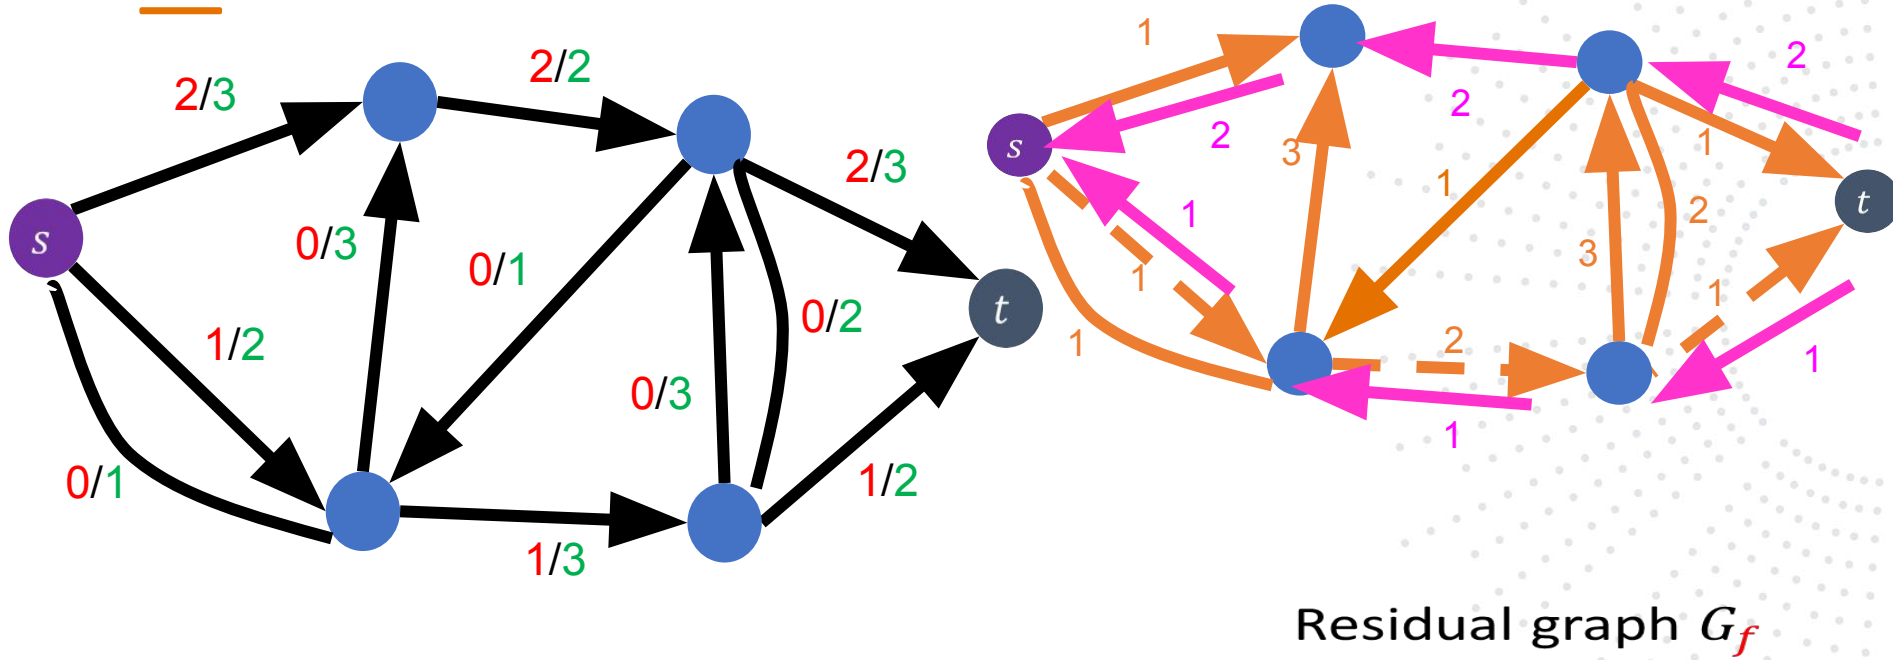


Ford-Fulkerson Algorithm- Example




Residual graph G_f


Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm- Example



Ford-Fulkerson Algorithm

Define an augmenting path to be an $s \rightarrow t$ path in the residual graph G_f (using edges of non-zero weight)

Ford-Fulkerson max-flow algorithm:

- Initialize $f(e) = 0$ for all $e \in E$
- Construct the residual network G_f
- While there is an augmenting path p in G_f :
 - Let $c = \min_{e \in p} c_f(e)$ ($c_f(e)$ is the weight of edge e in the residual network G_f)
 - Add c units of flow to G based on the augmenting path p
 - Update the residual network G_f for the updated flow

Initialization: $O(|E|)$


Construct residual network: $O(|E|)$

Finding augmenting path in residual network: $O(|E|)$ using BFS/DFS


Worst-Case Ford-Fulkerson Algorithm

- The **worst-case time complexity** of the **Ford-Fulkerson algorithm** can arise in specific scenarios where the algorithm has to process a large number of augmenting paths to reach the maximum flow
 - Path augmentation is slow
 - Edge capacities are small

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph

Increase flow by 1 unit




Residual graph

Worst-Case Ford-Fulkerson Algorithm




Flow graph




Residual graph

Worst-Case Ford-Fulkerson Algorithm



Flow graph

Increase flow by 1 unit



Residual graph

Ford-Fulkerson algorithm

- **Worst Case (Exponential Time Complexity):**
 - $O(E \cdot \text{Max Flow})$, which can be exponential in some cases
- This happens when the algorithm uses **Depth-First Search (DFS)** to find augmenting paths, and capacities lead to tiny flow increments.
- For example, if each augmenting path adds only 1 unit of flow and the maximum flow is F , there can be F iterations. If E edges are checked in each iteration, the time complexity becomes $O(E \cdot F)$

Edmonds-Karp Algorithm

- The **Edmonds-Karp algorithm** is a specific implementation of the **Ford-Fulkerson method** for solving the **maximum flow problem** in a flow network. It improves the Ford-Fulkerson method by using **Breadth-First Search (BFS)** to find the shortest augmenting path in terms of the number of edges

Edmonds-Karp Algorithm

Ford-Fulkerson max-flow algorithm:

- Initialize $f(e) = 0$ for all $e \in E$
- Construct the residual network G_f
- While there is an augmenting path in G_f , let p be the path with fewest hops:
 - Let $c = \min_{e \in E} c_f(e)$ ($c_f(e)$ is the weight of edge e in the residual network G_f)

How to find this?
Use breadth-first search (BFS)!

Edmonds-Karp = Ford-Fulkerson using
BFS to find augmenting path

Edmonds-Karp Algorithm

- **Edmonds-Karp algorithm** has a time complexity of $O(V \cdot E^2)$
 - BFS takes $O(V+E)$ time per iteration
 - Each augmenting path can increase the flow by a finite amount.
 - The maximum number of BFS iterations is $O(V \cdot E)$, because each edge is involved at most V times.
- Since each BFS takes $O(V+E)$, and BFS is called $O(V \cdot E)$ times, the total time complexity is:

$$O((V+E) \cdot V \cdot E) = O(V \cdot E^2)$$